


Q-Learning Convergence to Nash Equilibrium

Dashed green line = Nash equilibrium prediction

Hawk-Dove

Prisoner's Dilemma



Agent — Agent 1 — Agent 2